

# AIL 722: Assignment 2 Report

Anamitra Singha  
IIT Delhi

October 2024

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Model-Based Methods</b>	<b>2</b>
2.1	Policy Iteration . . . . .	2
2.1.1	Algorithm and Mathematical Explanation . . . . .	2
2.1.2	Implementation . . . . .	2
2.1.3	Visualization . . . . .	3
2.2	Value Iteration . . . . .	3
2.2.1	Algorithm and Mathematical Explanation . . . . .	3
2.2.2	Implementation . . . . .	3
2.2.3	Visualization . . . . .	4
<b>3</b>	<b>Model-Free Methods</b>	<b>4</b>
3.1	SARSA and Q-Learning . . . . .	4
3.1.1	Algorithm and Mathematical Explanation . . . . .	4
3.1.2	Implementation . . . . .	5
3.1.3	Results and Visualizations . . . . .	5
3.2	Evaluation and Analysis . . . . .	7
<b>4</b>	<b>Large-and-Continuous State-Space Environments</b>	<b>8</b>
4.1	LunarLander-v2 . . . . .	8
4.1.1	Algorithm and Neural Network Architecture . . . . .	8
4.1.2	Results and Visualization . . . . .	9

# 1 Introduction

This report presents the implementation of several reinforcement learning algorithms, applied to environments such as *TreasureHunt-v1*, *Taxi-v3*, *LunarLander-v2*, and *TreasureHunt-v2*. The report will cover the theoretical background, implementation details, and the performance of these algorithms, followed by visualizations and observations.

## 2 Model-Based Methods

### 2.1 Policy Iteration

#### 2.1.1 Algorithm and Mathematical Explanation

Policy Iteration alternates between policy evaluation and policy improvement. The goal is to find the optimal policy  $\pi^*$  by improving a policy  $\pi$  iteratively.

The Bellman equation for policy evaluation:

$$V^\pi(s) = \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^\pi(s')]$$

where:

- $V^\pi(s)$ : Value function for policy  $\pi$
- $P(s'|s, a)$ : Transition probability
- $R(s, a, s')$ : Reward function
- $\gamma$ : Discount factor

#### 2.1.2 Implementation

The Policy Iteration algorithm was implemented using Python. The environment used is *TreasureHunt-v1*, which consists of 400 states.

### 2.1.3 Visualization




Figure 1: Policy Visualization on the *TreasureHunt-v1* Grid (Arrows represent actions in each state).

Figure 2: Trajectory Visualization of Policy Iteration (GIF)

## 2.2 Value Iteration

### 2.2.1 Algorithm and Mathematical Explanation

Value Iteration is based on iteratively updating the value function using the Bellman optimality equation:

$$V(s) = \max_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V(s')]$$

### 2.2.2 Implementation

The Value Iteration algorithm was implemented, and it was applied to the *TreasureHunt-v1* environment.

### 2.2.3 Visualization




Figure 3: Value Function Visualization on *TreasureHunt-v1*.

### 3 Model-Free Methods

### 3.1 SARSA and Q-Learning

### 3.1.1 Algorithm and Mathematical Explanation

SARSA is an on-policy algorithm that updates the Q-values as:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

Q-Learning is an off-policy algorithm that updates the Q-values using:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

Where:

- $\alpha$ : Learning rate
- $\gamma$ : Discount factor
- $r$ : Reward

### 3.1.2 Implementation

Both SARSA and Q-Learning were implemented for *TreasureHunt-v1* and *Taxi-v3*. For *TreasureHunt-v1*, the learning rate  $\alpha = 0.3$ , discount factor  $\gamma = 0.95$ , and  $\epsilon = 0.4$ . For *Taxi-v3*,  $\alpha = 0.05$  and a decaying  $\epsilon$  was used.

### 3.1.3 Results and Visualizations




Figure 4: Reward vs Episodes for SARSA (*TreasureHunt-v1*).



Figure 5: Reward vs Episodes for Q-Learning (TreasureHunt-v1).

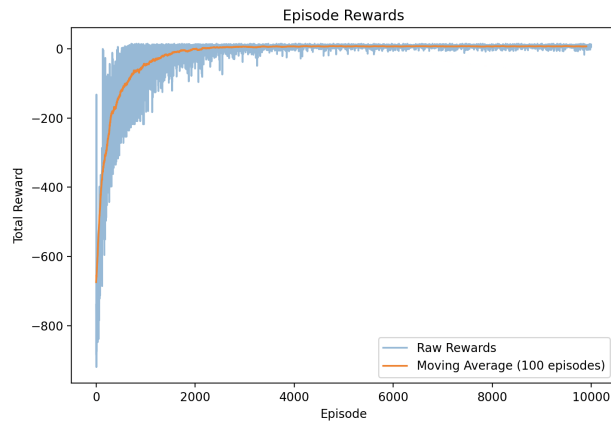


Figure 6: Epsilon vs Episodes for SARSA (Taxi-v3).




Figure 7: Epsilon vs Episodes for Q-Learning (Taxi-v3).

### 3.2 Evaluation and Analysis

For *TreasureHunt-v1* the time taken for convergence and the number of episodes were recorded, as shown in Table

Table 1: Comparison of SARSA and Q-Learning Convergence

Algorithm	Episodes for Convergence	CPU Time (s)
SARSA	13000	12
Q-Learning	10000	6

For *Taxi-v3*, the time taken for convergence and the number of episodes were recorded, as shown in Table

Table 2: Comparison of SARSA and Q-Learning Convergence

Algorithm	Episodes for Convergence	CPU Time (s)
SARSA	10000	10
Q-Learning	10000	6

## 4 Large-and-Continuous State-Space Environments

### 4.1 LunarLander-v2

#### 4.1.1 Algorithm and Neural Network Architecture

The Q-network used consists of four fully connected layers. The architecture is as follows:

- Input layer: The input state has a size of `state_size`.
- Hidden layer 1: Fully connected layer with 32 units, followed by a ReLU activation:

$$\text{fc1}(x) = \text{ReLU}(\mathbf{W}_1 x + \mathbf{b}_1) \quad \text{where} \quad \mathbf{W}_1 \in \mathbb{R}^{32 \times \text{state\_size}}$$

- Hidden layer 2: Fully connected layer with 64 units, followed by a ReLU activation:

$$\text{fc2}(x) = \text{ReLU}(\mathbf{W}_2 x + \mathbf{b}_2) \quad \text{where} \quad \mathbf{W}_2 \in \mathbb{R}^{64 \times 32}$$

- Hidden layer 3: Fully connected layer with 64 units, followed by a ReLU activation:

$$\text{fc3}(x) = \text{ReLU}(\mathbf{W}_3 x + \mathbf{b}_3) \quad \text{where} \quad \mathbf{W}_3 \in \mathbb{R}^{64 \times 64}$$

- Output layer: Fully connected layer with `action_size` outputs, representing Q-values for each action:

$$\text{fc4}(x) = \mathbf{W}_4 x + \mathbf{b}_4 \quad \text{where} \quad \mathbf{W}_4 \in \mathbb{R}^{\text{action\_size} \times 64}$$

Each layer uses the ReLU activation function except for the output layer, which outputs the raw Q-values for the given state-action pairs.



### 4.1.2 Results and Visualization




Figure 8: Reward Curve for *LunarLander-v2*.